

Lab 4: Threads

Nguyễn Thanh Hùng – 2540056

The 2D kernel maps thread (x, y) to pixel (y, x) . The notebook compares 1D and 2D blocks with the same number of threads and plots CPU/GPU speedup. Exercise 1: 16×16 permits four blocks and 1024 threads/SM; 8×8 permits only 512 threads with eight blocks; 32×32 exceeds 512 threads/block. Exercise 2: 128, 256, 512 and 1024 threads/block give 512, 1024, 1536 and 1024 resident threads respectively. Choose 512. Exercise 3: $\lceil 2000/512 \rceil 512 = 2048$ launched threads; 48 are inactive. These exercise limits are hypothetical; benchmarks use the actual GPU limits.

Kernel timing uses `time.time()` after JIT warm-up and CUDA synchronization, averaged over three GPU runs. Device allocation and host/device copies are excluded. CPU measurements are single runs. Speedup is relative to the stated CPU implementation. Actual measurements appear after running the notebook on Colab; no timings are assumed here.